

Automated Grading System

Winter Term Progress Report

CS 462

Wyman, Lucy

Hennig, Nathan

Gass, Andrew

Abstract

This document details the progress our group has made in the last three months on our Automated Grading System for OSU EECS. This includes elements we have begun working on and details of what code has already been written, as well as how we plan to build on these elements in the coming months. We also discuss problems and roadblocks we've come across, and other interesting aspects of development thus far. It also includes plans for the future of our project, and how we expect to complete the project.

I. PROJECT SUMMARY

The purpose of this project is to provide an automated grading system for EECS classes. This will allow teachers to quickly and asynchronously provide feedback to students, saving graders' time and improving feedback quality and response time for students.

Our system allows students to upload their homework to our server, which then runs the suite of tests uploaded by the teacher on that homework. Then, if the teacher wants students to be able to see feedback, the system will display which tests are or are not passing, along with hints for common errors students might run into. Both of these features can be enabled or disabled for a particular assignment by the teacher.

This project is particularly interesting in that we need to support tests in multiple languages, including Python, C/C++, Ruby, and JavaScript. Creating a system which can successfully run tests with limited interaction from teachers, students, or a system administrator has been the most challenging aspect of our project.

This project also has two interfaces, a command-line interface and web interface, both of which talk to the RESTful server to send and display data. This server then talks to the database to create, read, update, and delete appropriate data.

II. CURRENT STATE

We currently have the majority of the structure for our project erected, and are working on filling in details and trying to preempt any bugs that might come up as we go.

We are writing the command-line interface (CLI) in Python using a powerful module called 'cmd' that is part of the core Python library. The cmd module provides an extremely simple framework for implementing CLIs while being flexible enough for most purposes, including ours.

The CLI is complete except for final polishing and tweaking of output and error reporting. All command functions have been fully implemented. All usage information is printed dynamically from the Python dictionary we wrote that defines all commands for the CLI. Command output is printed neatly in columns. Using Python's basic logging module, we have implemented verbose debugging logs to provide debugging for the implemented commands. Implemented functions are running correctly and generating the expected output, normal and debug output alike.

One of the more important, yet conceptually simple, pieces of our project is the RESTful API listener server (herein referred to as the REST server) that will eventually live on the same server as our database. The listener server will catch HTML packets and query the database based on those packets and return the desired result.

The REST server is approximately ninety-five percent complete. All API calls are detected and responded to, but a few minor, non-core features (assignment tagging, for example) aren't working properly yet. Also, the API currently only returns output for view commands. Other commands such as insert and update just receive a code 200 OK response. Ideally all commands that modify the database will receive output about those modifications as their response.

The Web UI is also being written in Python, using Django. It is approximately 80% completed, with most of the structural components such as forms, sitemap, and API calls set up but not fully implements (i.e. the stub of a model exists without it being fully or correctly defined). The most interesting pages to date are the registration/login page, which display an example of how a form will look on the site and successfully communicate with the API to authenticate users (SESSIONS still to come!), and the home page, which displays the current courses and assignments for that individual student once they have logged in. We ended up scrapping the idea to use CAS and have decided to use our own methods of authentication. The benefit to this is that we have complete control over what information the students give us and are able to implement authentication as we go instead of waiting until the last minute, but the negative is that students have to individually register with our application separately instead of being automatically added as they are with Canvas for example. We decided this is worth the price of admission.

The other piece of authentication worth mentioning is that Professors must be manually added by an approved administrator to improve security, as we could not be certain that we would be able to create a secure way of identifying users as teachers or students in the time allotted. This is also a hassle, but it seemed better than the alternative of having students be able to authenticate as professors and change their grades, which would defeat the entire purpose of the application.

Currently, the majority of the work on the web interface will be filling in the stubs, creating the rest of the templates and urls, ensuring that forms and pages act as expected, and once the API is able to update the database as well as display it ensure that that functionality works through the GUI as well. This is entirely replicatory work, meaning that all new or unique features have been implemented successfully and now much be repeated and tweaked to fill out the rest of the application.

The last major piece of our project is the testing interface. This piece caused us some trouble at the beginning of the term, when we realized that we didn't have a firm idea of how our client wanted testing to be performed. After further discussion and research, we decided to implement an interface for each language we add testing capability for. The interfaces provide output in a standardized format, as laid out by the Test Anything Protocol (TAP). An example of TAP output follows:

```
1..4
# Basic tests
ok 1 one equals one
ok 2 one does not equal two
ok 3 a true statement
ok 4 testing works?
```

This was generated using the Python framework we implemented and a test that uses that framework to execute a few simple tests. In addition to expecting equality or inequality, the framework provides an expect error function, which can be passed a function handle that should throw an error. The structure for the testing functions and the Test Manager are somewhere between alpha and beta, and can listen to the RESTful API. We still need to expand the testing.

III. PROBLEMS

There were a number of problems with the CLI. The main problem was that we failed to consider the need to specifically design the command structure for the CLI, and failed to budget time to do so. The CLI design was completed around the end of week 3, although we are still fixing minor flaws as we find them, and the command structure is not fully consistent in all places (in other words, some commands use one argument pattern while similar commands use a different argument pattern).

Ideally, we will solve this problem by going back through the design and reworking it entirely. We may do that later depending on how much time we have to polish the interface.

The second problem with CLI was that after implementing about half the commands we noticed that code reuse was a problem. Many of the commands use code that similar to other commands, but not quite similar enough to be easily moved into a helper function.

We solved this problem by going through the design and making small changes that reduced syntactical variations between commands. Those changes allowed us to create a single set of execution functions that can process all commands based on their design in the Python command dictionary. This modular approach eliminated code reuse and makes it so that any changes to the command structure of the CLI will be easy to implement.

There have also been some unforeseen troubles with the Web UI. First is the URL and authentication issue. A few months ago we contacted EECS to ask about getting a subdomain for our application, and they notified us that even with the help of our mentor this wouldn't be possible until much closer to Expo. Because of this, we couldn't authorize the application with the Central Authentication Services, which we were planning on using for authentication. So, we decided to set up our own login system, initially thinking it would be temporary but then deciding that it was necessary to properly authenticate our users and would be production quality.

Then there were unexpected difficulties with setting up a basic login system to use in place of CAS for development. For a while I had a difficult time getting login to actually talk to my local database, but since that wasn't what we were planning on doing in production (the UI will talk to the database through an API, and it won't necessarily be local), I scrapped that and have decided to authenticate using the methods we'll be using in production, which do actually work. In retrospect this was a pretty minor and inconsequential roadblock, but at the time it was very frustrating.

As stated above in the status section, the testing section of our project gave us some issues after we realized that we hadn't adequately define requirements for that section during the design phase. We were able to solve this problem by researching various testing methods and looking at sites like CodeWars.com that focus on code testing. We eventually settled on a framework based on the CodeWars design for katas that generates Test Anything Protocol (TAP) output.

On the implementation side of the manager and testers, we used a multiprocessing approach to handle the expected load. An arbitrary number of testers may be launched with the manager, and which uses two threads to ensure that dispensing tests to testers never blocks responses to requests from either of the UIs. This should ensure expandability, stability, and responsiveness under load.

IV. PLANS

The CLI is essentially complete at this point, so we plan to spend time trying different ways of displaying output and fixing any bugs we discover as we continue testing.

For the REST API, we plan to continue testing and debugging it, as well as implementing the tags feature. We also want to finish implementing authentication through the database.

For the Web UI, all that is left to do is fill in the rest of the application, replicating and tweaking functionality we already have (ie. creating the rest of the forms and web pages, filling out the sitemap structure, and making all the correct API calls). We don't foresee any major roadblocks with this as the actual functionality has already been tested and is working, and now we are just filling in the stubs for the rest of the application.

The implementation of the testing framework should be fairly straightforward. We don't expect the initial issues to cause problems with hitting the beta release since we're now following a well-conceived design, and most of the work going forward will be linking with the database, as well as writing interfaces with more languages and deciding on the extent we plan to support multi-language tests.

V. CODE SNIPPETS

A. Web Interface (GUI)

Listing 1 is an example of querying the API for data to be displayed on a web page, parsing the data, and then rendering the page. This is part of the core functionality of the web interface part of the system, is just talking to the API, parsing that input/output, and then displaying it.

```
1 def index(request):
2     api_ip = settings.API_IP
3     user_data = {'student': ['hennign']}
4     userobj = requests.get(api_ip+'student/view', json=user_data)
5     user = userobj.json()
6     for i in range(len(user)):
7         user[i]['type'] = 'student'
8     courses = []
9     for course in user:
10        course_data = {'course-id': [course['course_id']]}
11        cobj = requests.get(api_ip+'course/view', json=course_data)
12        c = cobj.json()
13        aobj = requests.get(api_ip+'assignment/view')
14        c[0]['assignments'] = aobj.json()
15        courses.append(c[0])
16    return render_to_response('index.html', {'n': datetime.datetime.now(),
17        'user': user[0], 'courses': courses})
```

Listing 1: View function for the index page

The other core functionality is getting data from users to put in the database, which is done through forms like the one created in Listing 2.

```
1 class NewUser(forms.Form):
2     first_name = forms.CharField(label='First Name', required=True)
3     last_name = forms.CharField(label='Last Name', required=True)
4     ##TODO ask to re-enter onid
5     onid = forms.CharField(label='ONID ID Number', required=True)
6     usertype = forms.ChoiceField(choices=TYPES,
7         required=True, label='Are you a Professor, TA, or student?')
8     password = forms.CharField(label='Password',
9         required=True,
10        widget=forms.PasswordInput)
11    password2 = forms.CharField(label='Re-enter Password',
12        required=True,
13        widget=forms.PasswordInput)
```

Listing 2: Example of a form

B. Command Line Interface

The CLI is implemented using the module 'cmd', one of the core Python modules. Due to the robustness of said module, most of the CLI code is very straightforward. Here are a few samples of the code that weren't straightforward.

```

1 def print_response(self, command, subcommand, json):
2     data = json
3
4     cols = [x for x in sql_dict.sql[command][subcommand]['view_order']
5              if x in data[0]]
6     col_widths = [max([len(str(row[key])) for row in data] + \
7                       [len(str(key))]) + 4 for key in cols]
8
9
10    sort_order = [x
11                  for x in sql_dict.sql[command][subcommand]['sort_order']
12                   if x in data[0]]
13
14    print("|".join(str(val).center(col_widths[pos])
15                  for pos, val in enumerate(cols)))
16    print("|".join(str("="*(col_widths[pos]-2)).center(col_widths[pos])
17                  for pos, val in enumerate(cols)))
18
19    data.sort(key = lambda x: [x[key] for key in sort_order])
20
21    for row in data:
22        print("|".join(str(row[val]).center(col_widths[pos])
23                      for pos, val in enumerate(cols)))

```

Listing 3: CLI code responsible for printing the view commands results returned from the REST API

Listing 3 shows the CLI helper function `print_response`. This function takes a command, subcommand, and the json from a HTTP response and processes it into nice looking columns as seen in Figure 1. First the view and sort orders are checked against the data to see which columns are present in the data. Then a list of column widths is built using the maximum length of the values for each key or length of the key name, whichever is longer. After printing the column headers, the json data is sorted by the sort order keys and then each row is printed.

This piece of code is important, because this view output is what users will see most frequently. The output needs to be well formatted, both to impress the users and to improve the usability of the CLI.

```

>>> course view teacher=hennign
course_id | dept_name | course_num | name | term | year | teacher
=====
1 | cs | 161 | Introduction to Computer Science I | winter | 2016 | hennign
2 | ece | 271 | Digital Logic Design | spring | 2016 | hennign
3 | cs | 480 | Translators | winter | 2016 | hennign
>>> =

```

Fig. 1: Screenshot of CLI output from course view command showing column formatted data.

```

1 undoc_header = None
2 def print_topics(self, header, cmds, cmdlen, maxcol):
3     if header is not None:
4         cmd.Cmd.print_topics(self, header, cmds, cmdlen, maxcol)

```

Listing 4: CLI code responsible for hiding undocumented commands from help

By default, when the `help` command is run, undocumented commands are prefaced with the string in `undoc_header` and listed out. This behavior is undesirable, since the only undocumented commands are things like `EOF` and `exit`, where having help files would be awkward and confusing. Listing 4 demonstrates how we fix this issue.

The function `print_topics` is a built-in function of the `cmd` module from which the `AutoShell` class inherits and is used to print command lists for the various help categories. Here we override the function with our own copy which checks to see if the `header` is not equal to `None`. If this evaluates as true, then we call the original version of the function. If it is false, then we do nothing. This allows us to get the desired behavior without completely reimplementing the inherited function.

C. REST API

```
1 for table in tables:
2     query = """SELECT tc.constraint_name, tc.table_name, kcu.column_name,
3                   ccu.table_name AS foreign_table_name,
4                   ccu.column_name AS foreign_column_name
5                   FROM information_schema.table_constraints AS tc
6                   JOIN information_schema.key_column_usage AS kcu
7                   ON tc.constraint_name = kcu.constraint_name
8                   JOIN information_schema.constraint_column_usage AS ccu
9                   ON ccu.constraint_name = tc.constraint_name
10                  WHERE constraint_type = 'FOREIGN KEY'
11                  AND tc.table_name = %(tablename)s"""
12     cur.execute(query, table)
13     for row in cur.fetchall():
14         constraint, table, column, foreign_table, foreign_column = row
15         edges.append([row['table_name'], row['foreign_table_name'],
16                       row['column_name'], row['foreign_column_name']])
```

Listing 5: Generating a list of edges for tables in a PostgreSQL database from the information schema tables

Some of the most interesting code in the REST server is related to the way the API builds queries for the view commands. At first we considered writing each query by hand, but after careful consideration we decided that it would ultimately save time if the view queries were generated automatically. Relational databases can be thought of as graphs where each table is a node with edges created by foreign key constraints. PostgreSQL stores this data automatically in the `information_schema` tables. This data can later be extracted (see Listing 5) and used to represent a set of edges between tables.

The code in Listing 5 is important because it is the first step in building a graph that can be used to automatically generate joins for any given set of tables in the graph. Automatically generating joins this way allows easy modification of the view commands, thereby increasing the maintainability of the REST server.

```

1 elif 'multipart/form-data' in self.headers['Content-Type']:
2     form = cgi.FieldStorage(
3         fp=self.rfile,
4         headers=self.headers,
5         environ={'REQUEST_METHOD': 'POST',
6                 })
7
8     data = {}
9     for key in form.keys():
10         if key not in ['file', 'filepath']:
11             variable = str(key)
12             value = str(form.getvalue(variable))
13             self.logger.debug("value: {0}".format(value))
14
15             try:
16                 data[variable] = ast.literal_eval(value)
17             except:
18                 data[variable] = [str(value)]
19
20             if type(data[variable]) != type([]):
21                 data[variable] = [str(value)]
22
23     fileitem = form['file']
24
25     # Test if the file was uploaded
26     if fileitem.filename:
27         fn = os.path.basename(fileitem.filename)
28         open(os.path.normpath(sql['basedir'] + fn),
29             'wb').write(fileitem.file.read())

```

Listing 6: Processing a multipart/form-data HTTP request using cgi.FieldStorage and saving a file to disk

While the data from Content-Type application/json can easily be parsed, multipart/form-data is more complex. Luckily we determined that the cgi module was capable of automatically processing such a request, as shown in Listing 6. Once the data has been read into a FieldStorage object, getting the various variables was as simple as looping through the keys and pulling the values out. Listing 6 also shows how we saved the file data to disk for later use.

One of the more important tasks of the REST server is processing calls that include file data. One of the core requirements of our project is that users are able to submit files that are then tested. Without the code in Listing 6, any attempt to submit a file would fail.

D. Testing Framework

```
1 class dispenserThread (threading.Thread):
2     def __init__(self):
3         threading.Thread.__init__(self)
4     def run(self):
5         print('Dispenser thread running')
6         while True:
7             qlock.acquire()
8             while True:
9                 l = len(testQ)
10                if(l > 0):
11                    break
12            qlock.wait()
13            sub_ID = testQ.popleft()
14            qlock.release()
15            tester = -1
16            tlock.acquire()
17            while True:
18                try:
19                    tester = testers.index(0)
20                except Exception:
21                    pass
22                if(tester > -1):
23                    break
24            tlock.wait()
25            testers[tester] = int(sub_ID)
26            tlock.release()
27            s.connect('\0recvPort' + str(tester))
28            msg = '{"sub_ID":' + str(sub_ID) + '}'
29            s.send(msg.encode())
30            s.close()
```

Listing 7: Dispenser thread for my herald process. Wait, send.

The dispenser thread (Listing 7) works with the tester processes, checking for a pending test first, then checking for processes without a current test. Both are performed with Python condition blocking. In addition, upon assigning a test to a tester, the submission ID is stored in the tester's index. The main thread's 'matched' code:


```

1 while running:
2     readlist, writelist, exceptlist = select.select([c,k,r], [], [])
3     for avail in readlist:
4         if avail is c:
5             h = c.accept()[0]
6             msg = h.recv(256)
7             h.close()
8             dmsg = msg.decode()
9             sub_ID = json.loads(dmsg) ["sub_ID"]
10
11             #TESTING LINE
12             sub_ID = 83
13
14             qlock.acquire()
15             testQ.append(sub_ID)
16             qlock.notify()
17             qlock.release()
18         elif avail is k:
19             print("Herald recieved kill signal")
20             remv = k.accept()[0]
21             for j in range(int(sys.argv[1])):
22                 o = socket.socket(socket.AF_UNIX, socket.SOCK_STREAM);
23                 o.connect('\0killPort' + str(j))
24                 o.send('{"state":"die"}'.encode())
25                 o.close()
26                 k.close()
27             exit()
28         elif avail is r:
29             rec = r.accept()[0]
30             msg = rec.recv(256)
31             dmsg = msg.decode()
32             sub_ID = json.loads(dmsg) ["sub_ID"]
33             tlock.acquire()
34             try:
35                 idx = testers.index(int(sub_ID))
36                 testers[idx] = 0
37                 print('Submission evaluated.')
38             except ValueError:
39                 print("Error: a test was reported completed, but no record of
40
41
42

```

Listing 8: Main loop of my herald (manage) process.

The select tree (Listing 8) allows non-blocking port listening. If a new test is registered on the socket bound to 'c', then it is registered in a FIFO queue, and the thread condition is notified in case the dispenser was waiting on that. The 'r' socket will hear back from the testers and mark the tester process done, then notify of the freed process. The 'k' socket is a cleanup socket that closes out all the tester processes. Finally, the testing framework for Python is on the following page:

```

1 import operator
2
3 class test_suite:
4     def __init__(self):
5         self.testcount = 0
6         self.testsremain = 0
7
8     #Helper functions
9     def ok(self,message):
10         self.testcount += 1
11         self.testsremain -= 1
12         if(self.testsremain < 0): print(
13             '# Exceeded declared test count for this "describe"!')
14         print('ok ' + str(self.testcount) + ' ' + message)
15     def notok(self,message):
16         self.testcount += 1
17         self.testsremain -= 1
18         if(self.testsremain < 0): print(
19             '# Exceeded declared test count for this "describe"!')
20         print('not ok ' + str(self.testcount) + ' ' + message)
21
22     #Testing functions
23     def assert_equals(self, actual, expected, message):
24         if(operator.eq(actual, expected)):
25             self.ok(message)
26         else:
27             self.notok(message)
28     def assert_not_equals(self, actual, unexpected, message):
29         if(operator.ne(actual, unexpected)):
30             self.ok(message)
31         else:
32             self.notok(message)
33     def expect_error(self, message, thunk):
34         try:
35             thunk()
36             self.notok(message)
37         except:
38             self.ok(message)
39     def expect(self, passed, message):
40         if(passed):
41             self.ok(message)
42         else:
43             self.notok(message)
44
45     #Grouping functions
46     def describe(self, message, tests):
47         print(str(self.testcount + 1) + '..' + str(self.testcount + tests))
48         print('# ' + message)
49         self.testsremain = tests
50     def it(self, message):
51         print('# ' + message)

```

Listing 9: Python testing framework. Used to generate TAP output from a test suite.

Here in the test_suite (Listing 8) are the implementations for the testing functions. They reference back to print functions for their results, and the comparison relies on Python's operator library, a functionality we may need to implement separately in other languages.

VI. SCREENSHOTS

Below are screenshots of relevant aspects of our system which are particularly visual, or can be clarified visually.

A. Web Interface (GUI)

The index page, being the home of the site and the first page is most people see, so it needs to be done right! Students should ideally be able to access most, if not all, other parts of the website from this page. For that reason we included a list of courses and their relevant assignments, and will add links to upload a submission for that assignment if it is not past due.

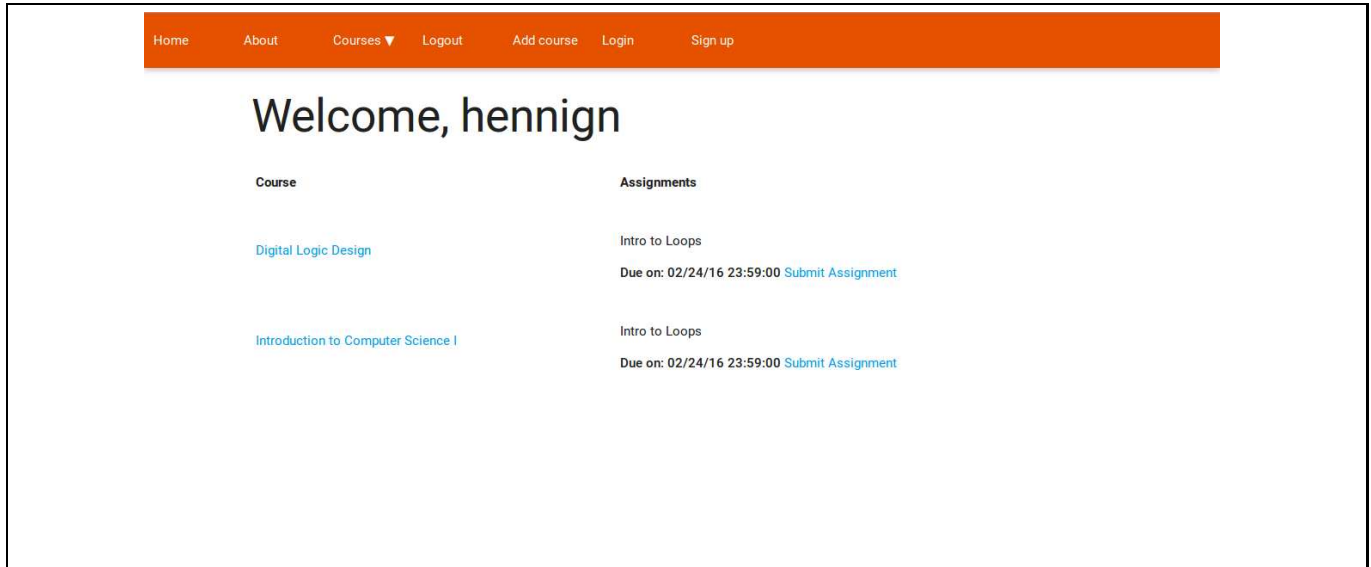


Fig. 2: Demo of the home page for the website

As stated previously, we also require the students register with our application individually, so needed a registration page as well.

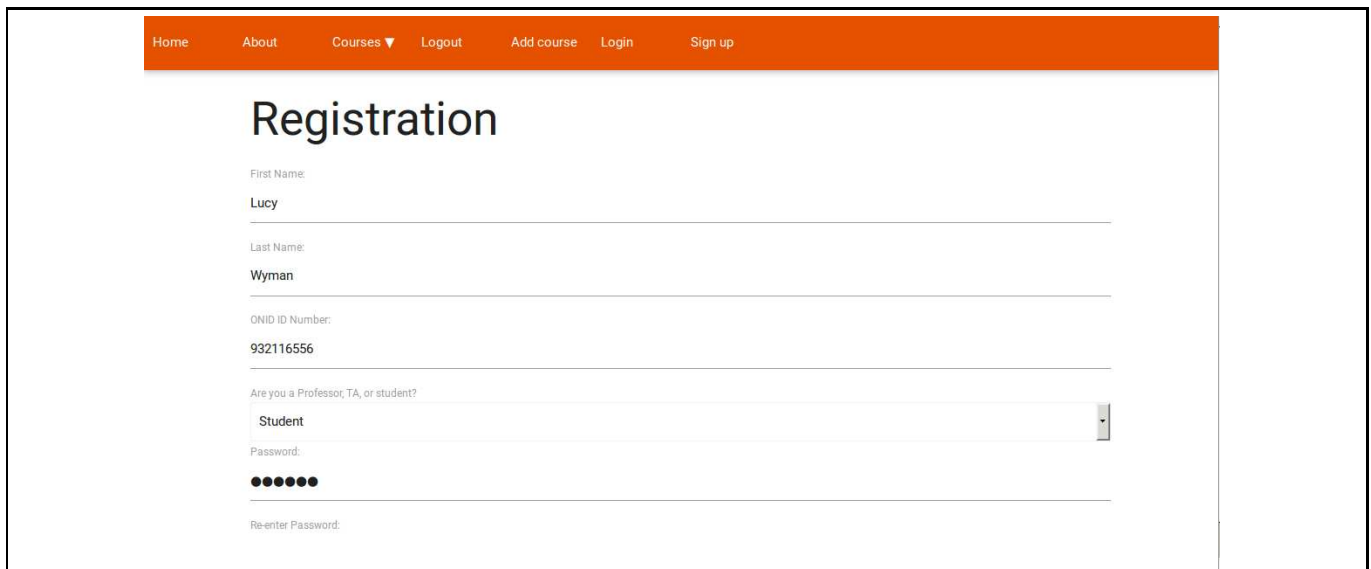


Fig. 3: Demo of the Registration page, where students must create an account

B. Command Line Interface

The nature of a command line interface is such that it doesn't really lend itself interesting screenshots, but in order to demonstrate how the CLI looks, we've included a few here.

```
Welcome to the AUTO Universal Testing Organizer (AUTO) shell.
Type help or ? to list commands
>>> ?

Documented commands (type help <topic>):
=====
assignment ce course grade group help student submission ta tag test
>>> help ce

**** Access denied
>>> level teacher
Now a teacher!
>>> help ce

NAME
    ce

SYNOPSIS
    ce link <ce-id=<value> test-id=<value>>
    ce delete <ce-id=<value>>
    ce update <ce-id=<value> [name=<value> text=<value>]
    ce view [assignment-id=<value> ce-id=<value> course-id=<value> name=<value>
            test-id=<value> version=<value>]
    ce add <name=<value> text=<value>>

DESCRIPTION
    ce link      Link selected common error(s) to selected test(s).
    ce delete    Delete selected common error(s).
    ce update    Update a common error with chosen values.
    ce view      View all common errors [optionally filtered by key-value pairs].
    ce add       Add common error with specified values.

OPTIONS
    assignment-id  Number that identifies a particular assignment. Displayed using assignment
                    view.
    ce-id          Integer number that identifies a common error. Displayed using ce view.
    course-id      Integer number that identifies a course. Displayed using course view.
    name           Can be any quoted string
    test-id        Number that identifies a particular test. Displayed using test view.
    text           Filepath to file, or a string of text, explaining the error.
    version        Version number for assignment. Defaults to highest (most recent) available.

>>> _
```

Fig. 4: Demonstrating opening message of the CLI and the output of 'help ce'.

Figure 4 demonstrates a few interesting aspects of the CLI. The first is the welcome message, which welcomes the user and explains how to use the help command to see other commands. Next we use the 'level' command to change access levels while also using 'help ce' command to show how a help file looks different depending on your access level.

Figure 5 shows a test run of the `course` command with verbose on. Verbose is a testing command that will most likely either be removed from the released version or restricted to teachers. The CLI has been written using Python's logging module to insert debug messages. The verbose command allows the user to change the logging level from warning to debug and back.

As can be seen from the debug messages above, the attempt to run the `course` command worked beautifully. All the arguments were parsed correctly and the API call was made correctly. Unfortunately, the REST server doesn't implement the `course` command yet, so it returned a 501 'Unsupported method' error as its response.

```

>>> level teacher
Now a teacher!
>>> verbose on
Verbose: ON
>>> course add name="Computer Architecture" dept=cs num=472 term=spring year=2017
2016-03-13 22:13:52.680 - command_exec - DEBUG: START. Args='add name="Computer Architecture" dept=cs num=472 term=spring year=2017'
2016-03-13 22:13:52.680 - command_exec - DEBUG: Args split. Args=['add', 'name="Computer Architecture"', 'dept=cs', 'num=472', 'term=spring', 'year=2017']
2016-03-13 22:13:52.680 - command_access - DEBUG: Checking access levels.
2016-03-13 22:13:52.680 - command_access - DEBUG: User level is teacher.
2016-03-13 22:13:52.680 - command_access - DEBUG: Access for teacher is True
2016-03-13 22:13:52.680 - command_access - DEBUG: Access granted.
2016-03-13 22:13:52.680 - command_exec - DEBUG: Entering ADD mode
2016-03-13 22:13:52.680 - command_data - DEBUG: Entering argument processing.
2016-03-13 22:13:52.680 - command_data - DEBUG: Valid keys include ['name'], [], ['dept', 'num', 'term', 'year'].
2016-03-13 22:13:52.680 - command_data - DEBUG: data is '<'term': ['spring'], 'num': ['472'], 'name': ['Computer Architecture'], 'year': ['2017'], 'dept': ['cs']>'
2016-03-13 22:13:52.680 - command_data - DEBUG: Verifying that required options are present.
2016-03-13 22:13:52.680 - command_request - DEBUG: url is 'http://127.0.0.1:8000/course/add/'
2016-03-13 22:13:52.680 - command_request - DEBUG: POSTing
send: b'POST /course/add/ HTTP/1.1\r\nHost: 127.0.0.1:8000\r\nAccept-Encoding: gzip, deflate\r\nContent-Length: 105\r\nConnection: keep-alive\r\nContent-Type: application/json\r\n\r\n{"term": ["spring"], "num": ["472"], "name": ["Computer Architecture"], "year": ["2017"], "dept": ["cs"]}'
reply: 'HTTP/1.0 200 OK\r\n'
header: Server header: Date header: Content-type
No response
2016-03-13 22:13:52.700 - command_exec - DEBUG: END
>>> _

```

Fig. 5: Demonstration of the course command with debugging active.

VII. CONCLUSION

We are pleased to report that our beta version is essentially complete, just as we planned. All that is left for the core of our project is to properly integrate authentication. Right now there are a few places that should be checking for access rights that aren't, but we expect implementation to consist of little more than a few short functions. Otherwise we expect to be focused on polishing code and adding a few bonus features as schedules permit.